




USC Viterbi
School of Engineering

Lucene, Solr and HW#4






What is Lucene

- **High performance, scalable, full-text search library**
- **Basically one adds documents to Lucene which creates an inverted index**
 - A document is defined as a collection of fields
 - Lucene includes flexible text analysis including tokenizers and filters
- **Lucene does NOT include crawlers or explicit searching of documents**
 - Searching the index must be implemented elsewhere
- **100% Java, no dependencies**
- **Offered by the Apache Software Foundation**

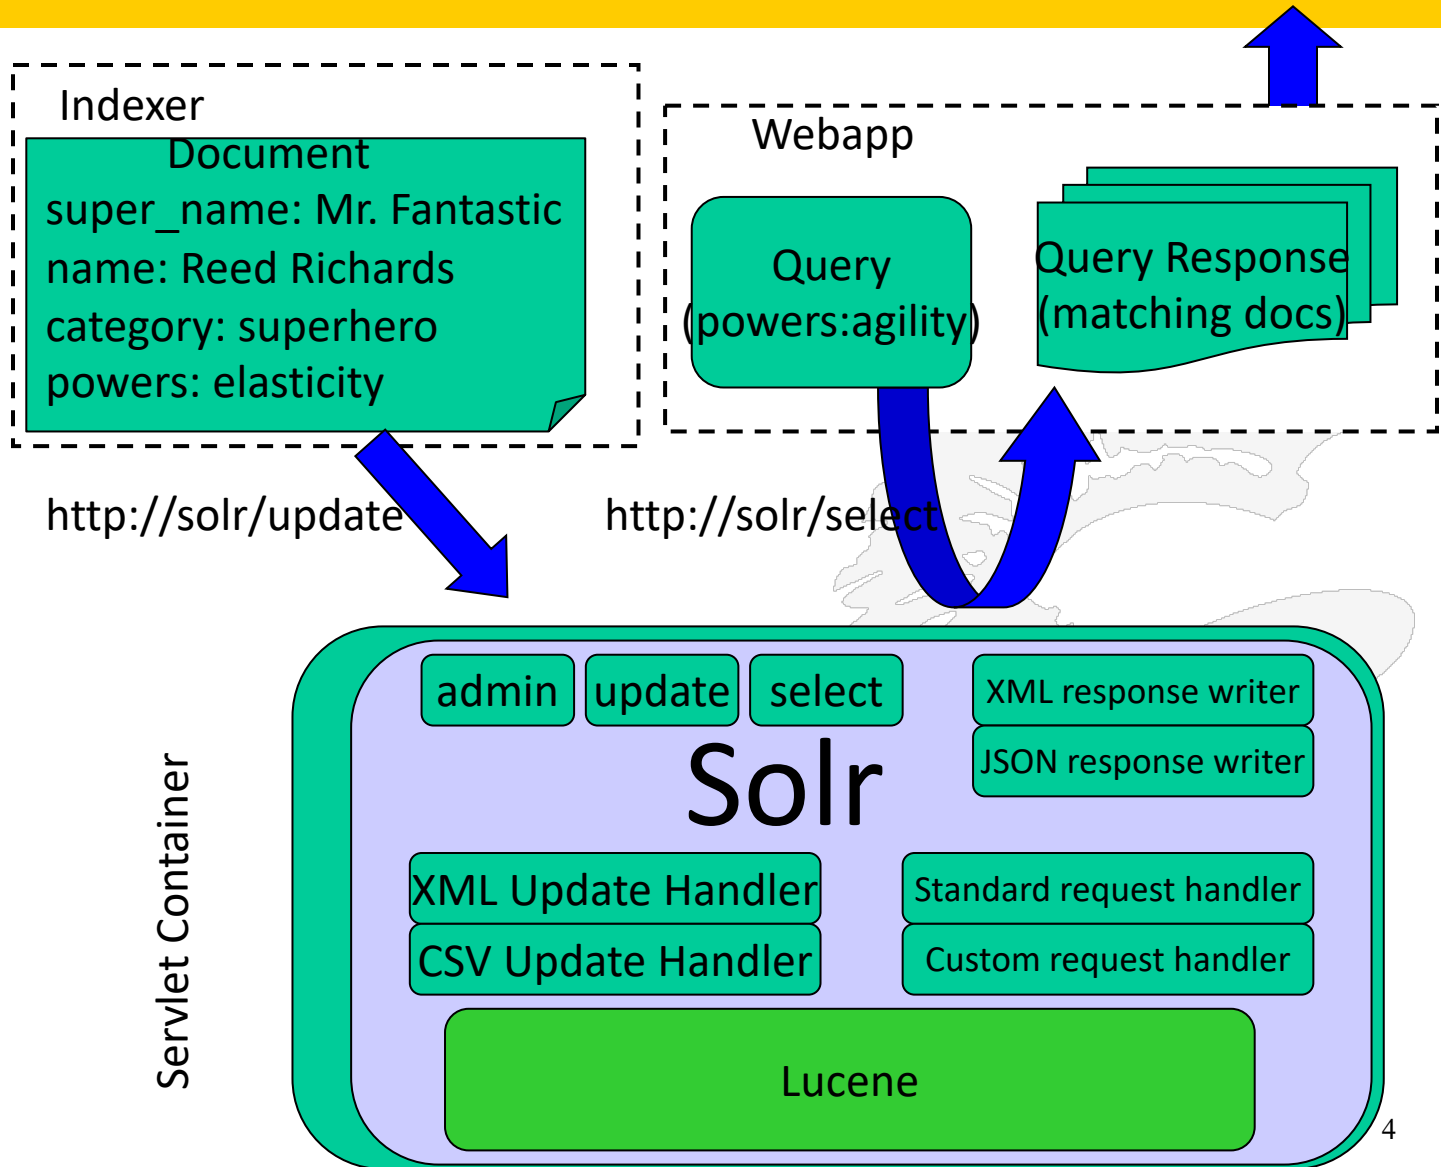


What is Solr

- **A full text search server based on Lucene**
- **Communicates with Lucene via XML/HTTP, JSON Interfaces**
- **Flexible data schema to define types and fields**
- **Output includes highlighting matches**
- **Configurable advanced caching**
- **Index Replication**
- **Extensible Open Architecture, Plugins**
- **Web Administration Interface**
- **Written in Java5, deployable as a WAR**

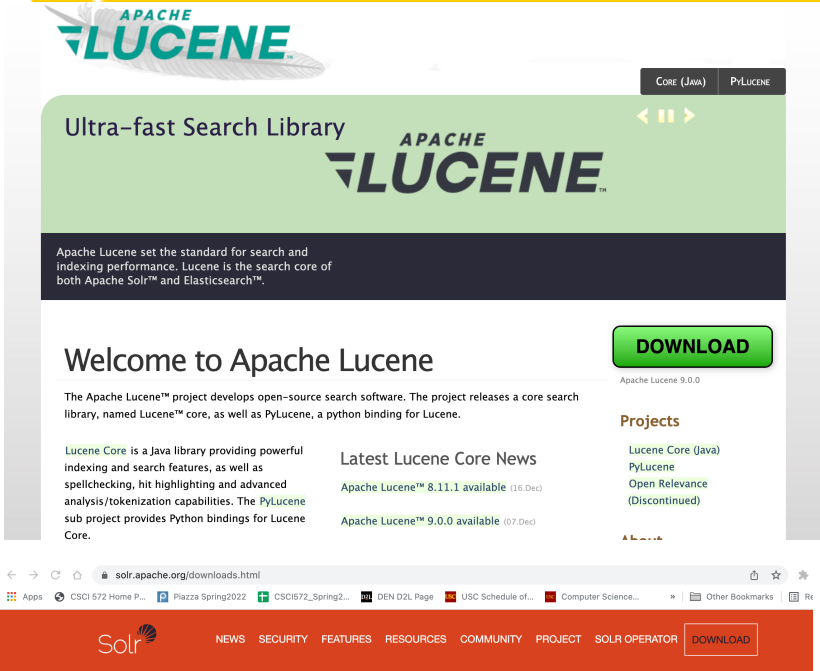


Top-Level Overview of Lucene/Solr





Downloading Solr



Lucene
home page

Solr
home page
Download
Solr

Solr Downloads

Official releases are usually created when the **developers** feel there are sufficient changes, improvements and bug fixes to warrant a release. Due to the voluntary nature of Solr, no releases are scheduled in advance.

Solr 8.11.1

Solr 8.11.1 is the most recent Apache Solr release.

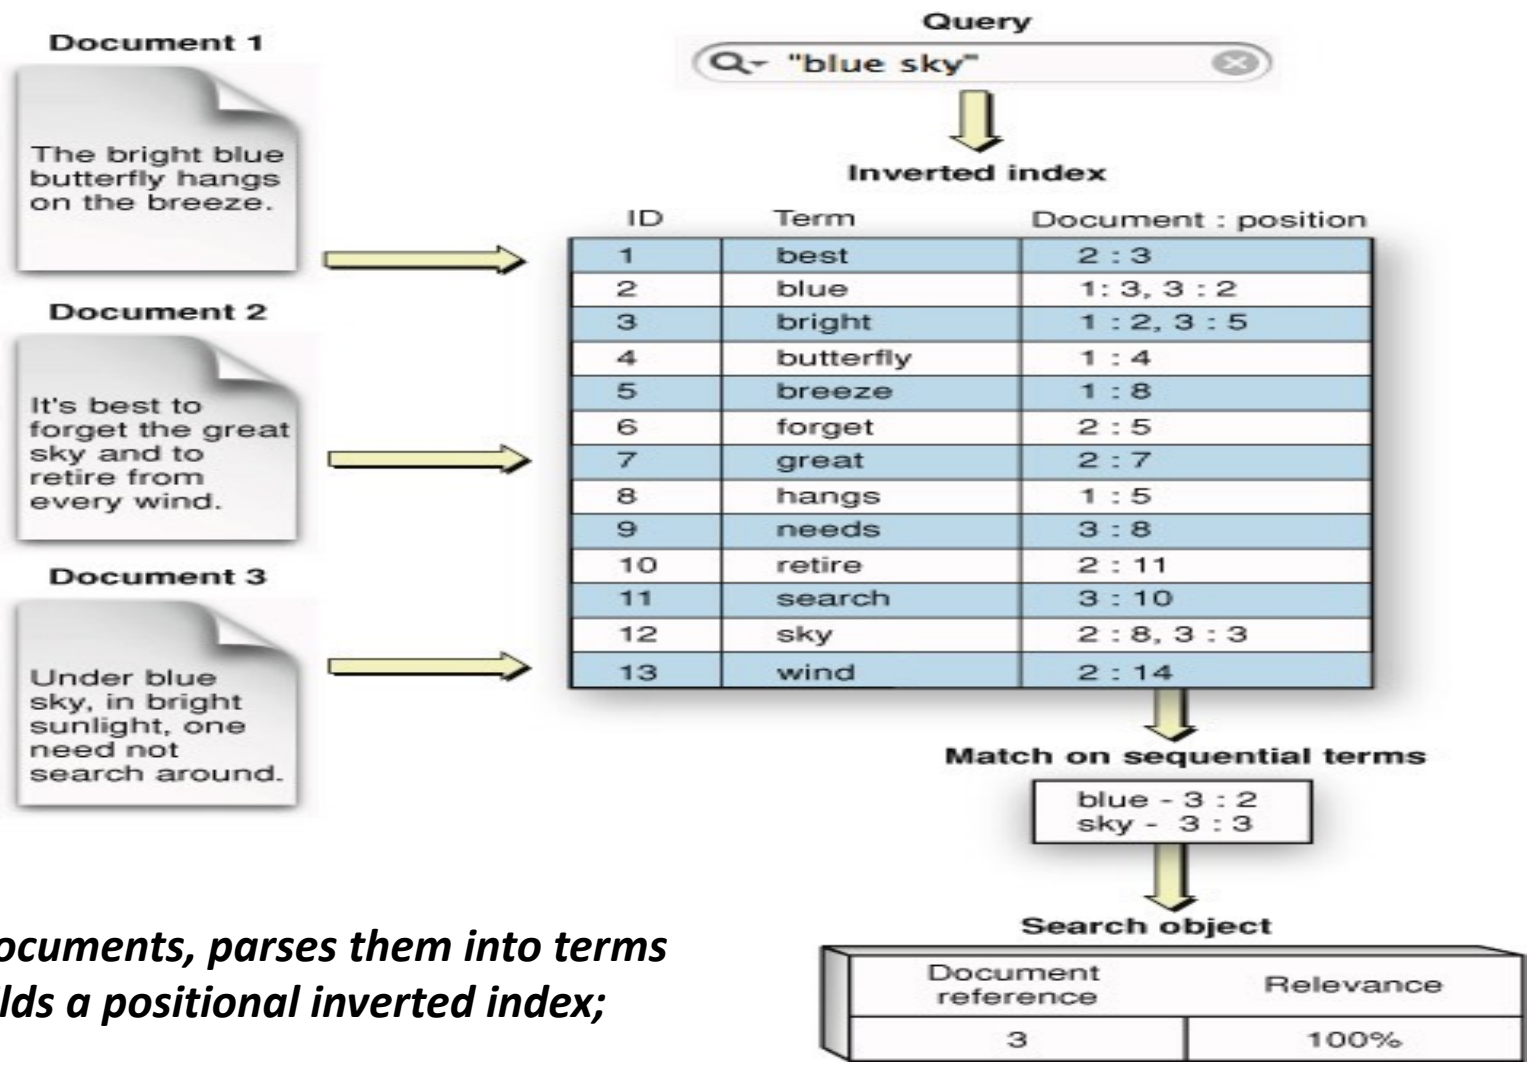
- Source release: [solr-8.11.1-src.tgz](#) [PGP] [SHA512]
- Binary releases: [solr-8.11.1.tgz](#) [PGP] [SHA512] / [solr-8.11.1.zip](#) [PGP] [SHA512]
- Docker: [solr:8.11.1](#)
- [Change log](#)

Solr 7.7.3

Solr 7.7.3 is the last release in the 7.x series.

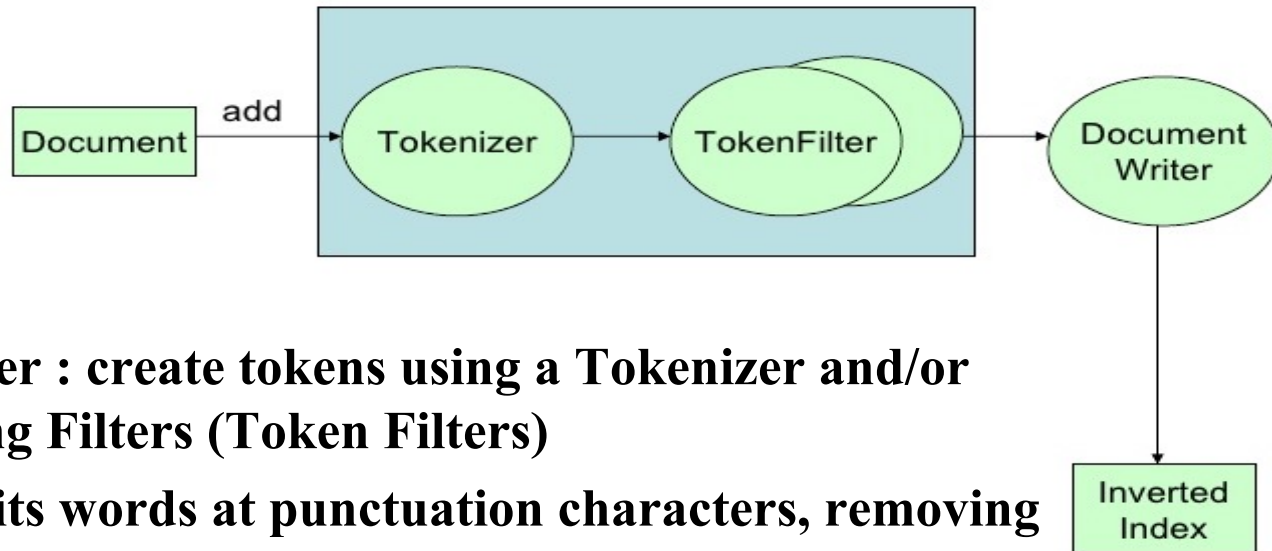
***You only need to download Solr
Due to security extensions in version 8.11
We recommend you use version 7.7.3***

Lucene Internals – Positional Inverted Index



*takes documents, parses them into terms
and builds a positional inverted index;*

Lucene Indexing Pipeline Analyzer



- **Analyzer : create tokens using a Tokenizer and/or applying Filters (Token Filters)**
 - Splits words at punctuation characters, removing punctuation. However, a dot that's not followed by whitespace is considered part of a token.
 - Splits words at hyphens, unless there's a number in the token, in which case the whole token is interpreted as a product number and is not split.
 - Recognizes email addresses and internet hostnames as one token.



Two Lucene Scoring Concepts

TF – IDF & cosine similarity

Lucene scores using a combination of TF-IDF and vector closeness

$$w_{x,y} = tf_{x,y} \times \log \left(\frac{N}{df_x} \right)$$

TF-IDF

Term x within document y

$tf_{x,y}$ = frequency of x in y

df_x = number of documents containing x

N = total number of documents

TF - IDF = Term Frequency X Inverse Document Frequency

$\text{cosine-similarity}(\text{query_vector}, \text{document_vector}) = V(q) * V(d) / |V(q)| * |V(d)|$
where $V(q) * V(d)$ is the dot product of the weighted vectors and $|V(q)|$, $|V(d)|$ are the Euclidean norms of the vectors (square root of the sum of squares)

for details see

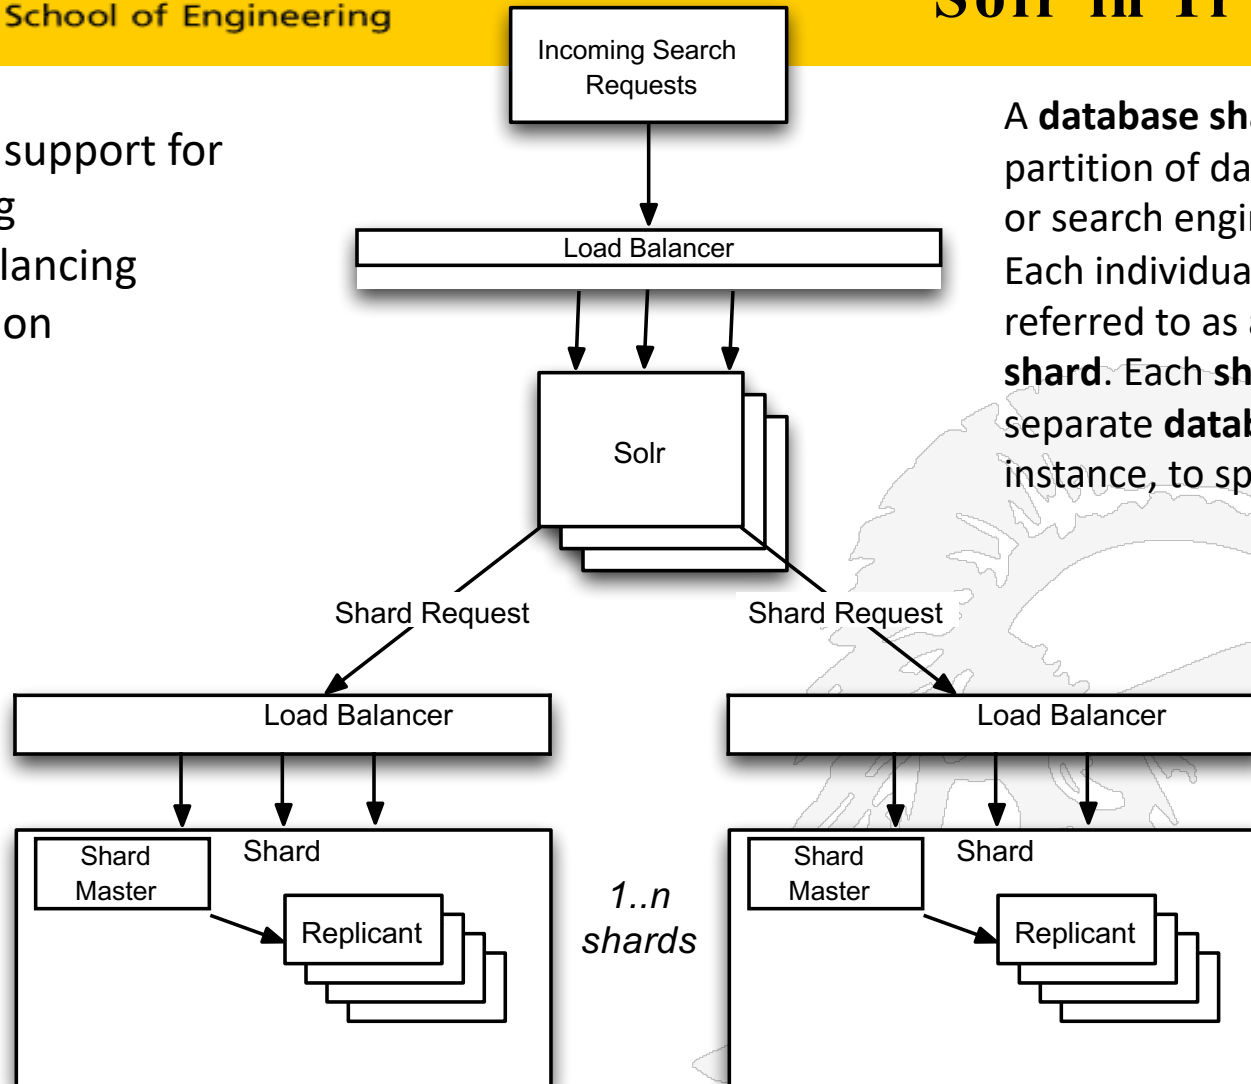
https://lucene.apache.org/core/4_0_0/core/org/apache/lucene/search/similarities/TFIDFSimilarity.html



Solr in Production

Includes support for

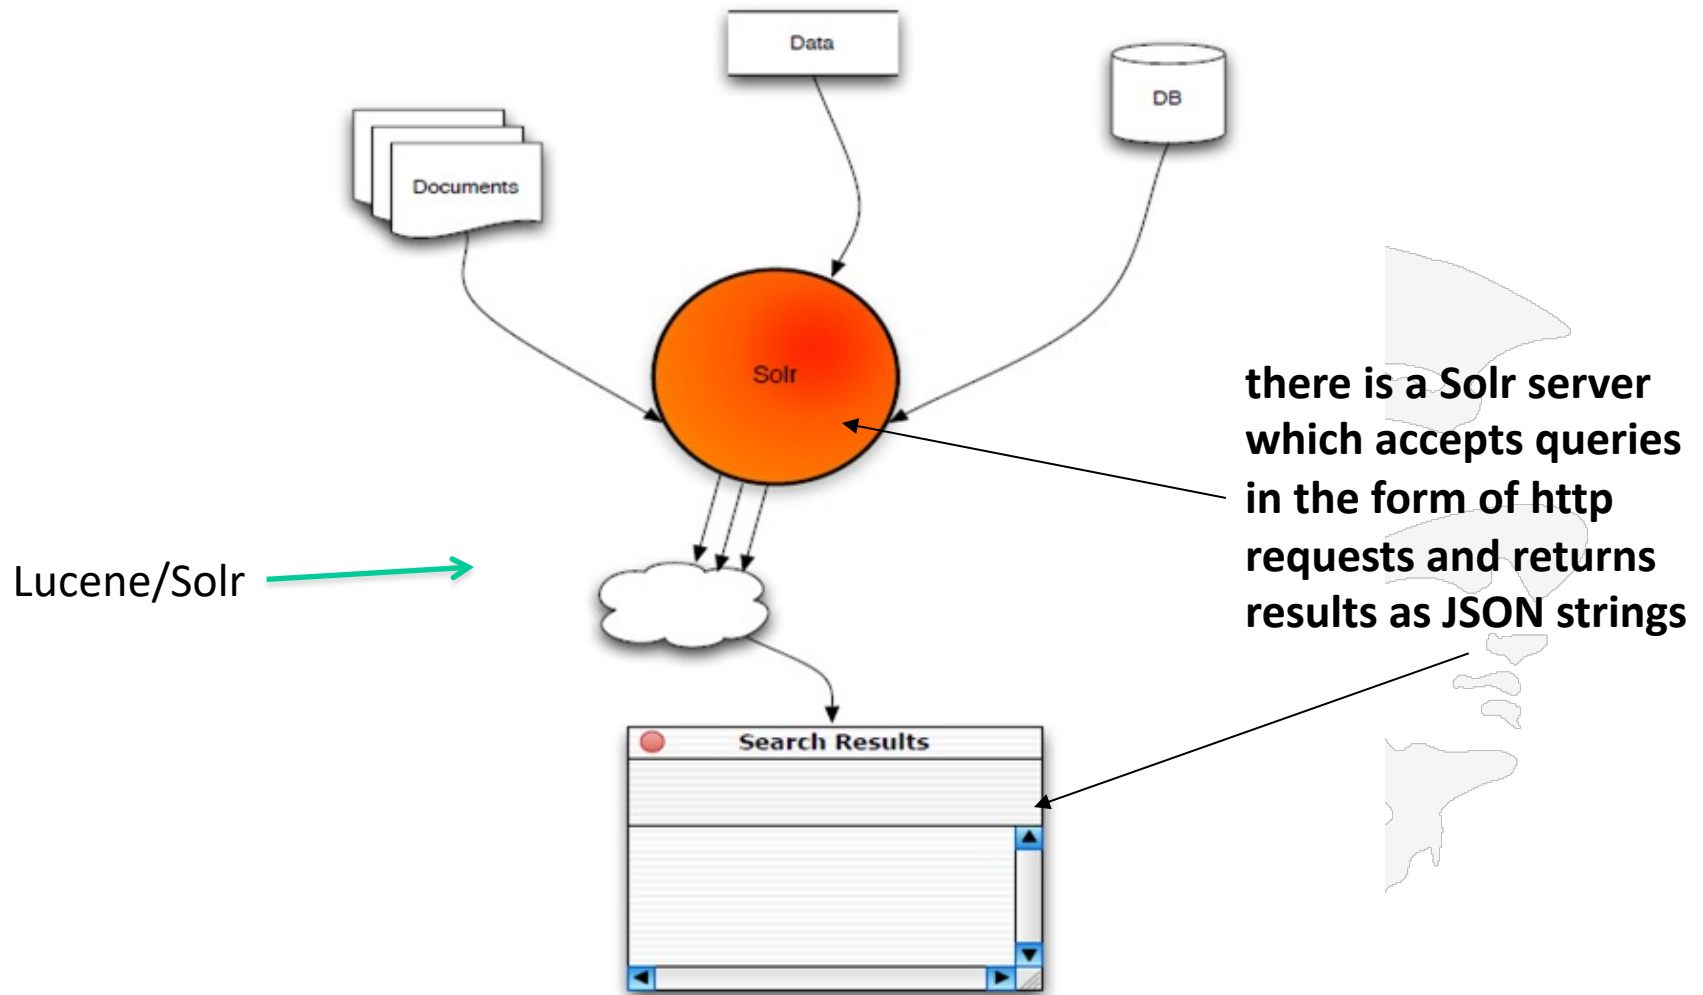
- Sharding
- Load balancing
- replication



A **database shard** is a horizontal partition of data in a **database** or search engine. Each individual partition is referred to as a **shard**. Each **shard** is held on a separate **database** server instance, to spread the load.



High Level Overview





Solr Request and Result

http://localhost:8983/solr/select?q=video&start=0&rows=2&fl=name,price

<response><responseHeader><status>0</status>

<QTime>1</QTime></responseHeader>

<result numFound="16173" start="0">

<doc>

<str name="name">Apple 60 GB iPod with Video</str>

<float name="price">399.0</float>

</doc>

<doc>

<str name="name">ASUS Extreme N7800GTX/2DHTV</str>

<float name="price">479.95</float>

</doc>

</result>

</response>

- status, always 0 in a successful response
- QTime, the server-side query time in milliseconds
- numFound, the total number of documents matching the query
- start, the offset into the ordered list of results
- field types in <doc> include str, boolean, int, long, float, double, date, lst, arr
 - lst is a named list <lst><int name="foo">33</int><int name="bar">42</int></lst>
 - arr is an array <arr><int>33</int><int>42</int></arr>
- multivalued fields are returned in an <arr> element.



How to Start Solr

Complete installation instructions can be found at

https://solr.apache.org/guide/8_11/installing-solr.html

Once it is installed:

1. Start Solr


```
java -jar start.jar
```

2. Index your XML data

```
java -jar post.jar *.xml
```

3. Search <http://localhost:8983/solr>

localhost indicates the Solr server is running locally on port 8983





USC Viterbi
School of Engineering

Querying Data and XML Response

- HTTP GET or POST with parameters are used to specify queries
- E.g. here are 4 sample queries, some with various parameters

`http://solr/select?q=electronics`

`http://solr/select?q=electronics&sort=price+desc`

`http://solr/select?q=electronics&rows=50&start=50`

`http://solr/select?q=electronics&fl=name+price`

price and desc are fields

limit results to fields:
name and price

```
<response>
  <lst name="responseHeader">
    <int name="status">0</int>
    <int name="QTime">1</int>
  </lst>
  <result name="response" numFound="14" start="0">
    <doc>
      <arr name="cat">
        <str>electronics</str>
        <str>connector</str>
      </arr>
```

...



Querying Data: Results

- The standard response format is JSON

```
{
  "responseHeader":{
    "zkConnected":true,
    "status":0,
    "QTime":8,
    "params":{
      "q":"foundation"}},
  "response":{"numFound":4,"start":0,"maxScore":2.7879646,"docs":[
    {
      "id":"0553293354",
      "cat":["book"],
      "name":"Foundation",
      "price":7.99,
      "price_c":"7.99,USD",
      "inStock":true,
      "author":"Isaac Asimov",
      "author_s":"Isaac Asimov",
      "series_t":"Foundation Novels",
      "sequence_i":1,
      "genre_s":"scifi",
      "_version_":1574100232473411586,
      "price_c____l_ns":799}]
  }}
```





Query Types Supported by Solr

- **Single and multi-term queries**
 - ex `fieldname:value` or `title: software engineer`
- **+, -, AND, OR, NOT operators are supported**
 - ex. `title: (software AND engineer)`
- **Range queries on date or numeric fields,**
 - ex: `timestamp: [* TO NOW]` or `price: [1 TO 100]`
- **Boost queries:**
 - e.g. `title:Engineer ^1.5 OR text:Engineer`
- **Fuzzy search :** is a search for words that are similar in spelling
 - e.g. `roam~0.8 => noam`
- **Proximity Search :** with a sloppy phrase query. The closer together the two terms appear, higher the score.
 - ex `“apache lucene”~20` : will look for all documents where “apache” word occurs within 20 words of “lucene”



Solr/Lucene is Used WorldWide

- **Search Engine**
 - **Yandex.ru, DuckDuckGo.com**
- **Newspaper**
 - **Guardian.co.uk**
- **Music/Movies**
 - **Apple.com, Netflix.com**
- **Events**
 - **Stubhub.com, Eventbrite.com**
- **Cloud Log Management**
 - **Loggly.com**
- **Others**
 - **Whitehouse.gov**
- **Jobs**
 - **Indeed.com, Simplyhired.com, Naukri.com**
- **Auto**
 - **AOL.com**
- **Travel**
 - **Cleartrip.com**
- **Social Network**
 - **Twitter.com, LinkedIn.com, mylife.com**

[Apache Solr Popularity](#)


Usage statistics for *Apache Solr Popularity*

This page provides information about the usage of the *Apache Solr Popularity* project, including summaries across all versions and details for each release. For each week beginning on the given date the figures show the number of sites that reported they are using a given version of the project.

These statistics are incomplete; only Drupal websites using the [Update Status](#) module are included in the data. This module has been included with the download of Drupal since version 6.x so the data does not include older sites. [Read more information about how these statistics are calculated.](#)

[Apache Solr Popularity project page](#)
[Usage statistics for all projects](#)

Weekly project usage



Tracking the popularity of Solr
 According to DB-Engines
 Elasticsearch is ranked number one,
 and Solr is ranked number three



Solr Includes Spell Checking

(you will implement this for homework #5)

- **Not enabled by default, see example config to wire it in**
- https://lucene.apache.org/solr/guide/7_7/spell-checking.html
- **Uses file or index-based dictionaries for spell correction**
- **Supports pluggable distance algorithms:**
 - **Levenstein alg:** https://en.wikipedia.org/wiki/Levenshtein_distance
 - **JaroWinkler alg:** ,
https://en.wikipedia.org/wiki/Jaro%E2%80%93Winkler_distance
- <https://cwiki.apache.org/confluence/display/solr/SpellCheckComponent> is a full discussion of the spellchecking abilities of Solr

Solr Includes Autosuggestion

(you will implement this for homework #5)

Enter your keywords:

teach

Did you mean: **teaching**

teach	17
teachers	2
teacher	1
teach book	15
teach world	11
teach wide	11
teach teaching	9
teach computer	9

Find dinn|

dinner

dinner restaurant

dinner and drinks

dinner cruise

dinner and dancing

dinner date

dinner theater

dinner show

dinner buffet

dinner and live jazz

https://lucene.apache.org/solr/guide/7_7/suggester.html



HW#4 Exercise

There are several components to this homework; you will possibly not need all of them;

- 1. Installing Ubuntu on Windows**
- 2. Installing Solr**
- 3. Using Solr to Index a web site**
- 4. The actual exercise - comparing ranking algorithms**
- 5. What to submit**
- 6. Grading guidelines**





What you will need to do

SUMMARY (1 of 2)

1. If you are using Windows, you will have to install Ubuntu on your machine
 - if you have a Mac, skip to step 2
 - Solr requires a Unix platform
 - first you must download and install Oracle's VirtualBox, free software that will permit you to install Ubuntu on your Windows device
 - second you will download and install Ubuntu within Virtualbox
2. Download and install the current release of Solr
 - there is a Quick Start Guide to help you at https://lucene.apache.org/solr/guide/7_7/solr-tutorial.html
 - the Solr server should be started; to check that it is running look at <http://localhost:8983/solr>
3. Download the reference news website you are responsible for
 - 3 news websites are located on Google drive and are accessible using
 - <https://drive.google.com/drive/folders/14qtJywCRk00wEbmt-Pv3efyPIBqwcoZr>



What you will need to do

SUMMARY (2 of 2)

4. Import the news website web pages you are responsible for into Solr
 - There are only three new websites, use last digits of your USC ID to determine which one you are responsible for
5. Using your laptop install or use an existing web server to create a website
 - Macs (I believe) have Apache pre-installed
 - in your website create a web page that looks like a Google query box
 - in your website write a program that takes a query and sends it to Solr
 - in your website write a program that accepts the results from Solr and displays them as a web page
6. Run the set of 9 queries at the end of SolrExercise.pdf
 - save the top ten results for each of the 9 queries
7. Use the NetworkX library, the downloaded web pages and Jsoup library to create a network graph of the downloaded web pages
 - use the PageRank function from NetworkX to create a file, pagerankFile.txt that contains individual page ranks for every downloaded web page
 - install the pagerankFile in Solr and direct Solr to use it on queries
8. Run the set of 9 queries at the end of SolrExercise.pdf
 - save the top ten results for each of the queries
9. compare the two sets of results and describe any overlap



Step 1: Ubuntu with VirtualBox

- **VirtualBox is an open source, freely available Windows application (it also runs on other platforms) that lets you run multiple operating systems on your single machine**
 - E.g. run Windows on a Mac, run Linux on Windows
 - Major supported operating systems include: Windows NT 4.0, Windows 2000, Windows 8, Windows 10, DOS Windows 3.x, Linux, Solaris, FreeBSD, OpenBSD
- **Ubuntu is a Linux-based operating system distributed on personal computers, smartphones and network servers. It uses Unity as its default desktop environment**
- **Solr requires a Unix environment to run, so step 1 is required if you plan to use your Windows laptop**



Step 1: Setting Up Ubuntu with VirtualBox

1. **Download the free version of VirtualBox for Windows machines**
 - **Instructions can be found here**
<http://csci572.com/2022Fall/hw4/UbuntuVirtualBoxFinal.pdf>
2. **Download the Ubuntu 64-bit version**
3. **Run VirtualBox and select your Ubuntu version as the New Application**
4. **Set various parameters**
5. **Install Ubuntu and you should be ready to run**

Your Ubuntu Desktop





Step 2: Installing Solr


- **Solr is an open source enterprise search server based on the Lucene Java search library**
- **Instructions for downloading and installing Solr can be found here**
 - https://solr.apache.org/guide/7_7/installing-solr.html
- **Lucene is a fast, high performance, scalable search/information retrieval library**
 - Initially developed by Doug Cutting (Also author of Hadoop)
 - it provides for Indexing and Searching of documents
 - produces an Inverted Index of documents
 - Provides advanced Search options like synonyms, stopwords, based on similarity, proximity.



USC Viterbi
School of Engineering

Downloading Your Data Set of Web Pages

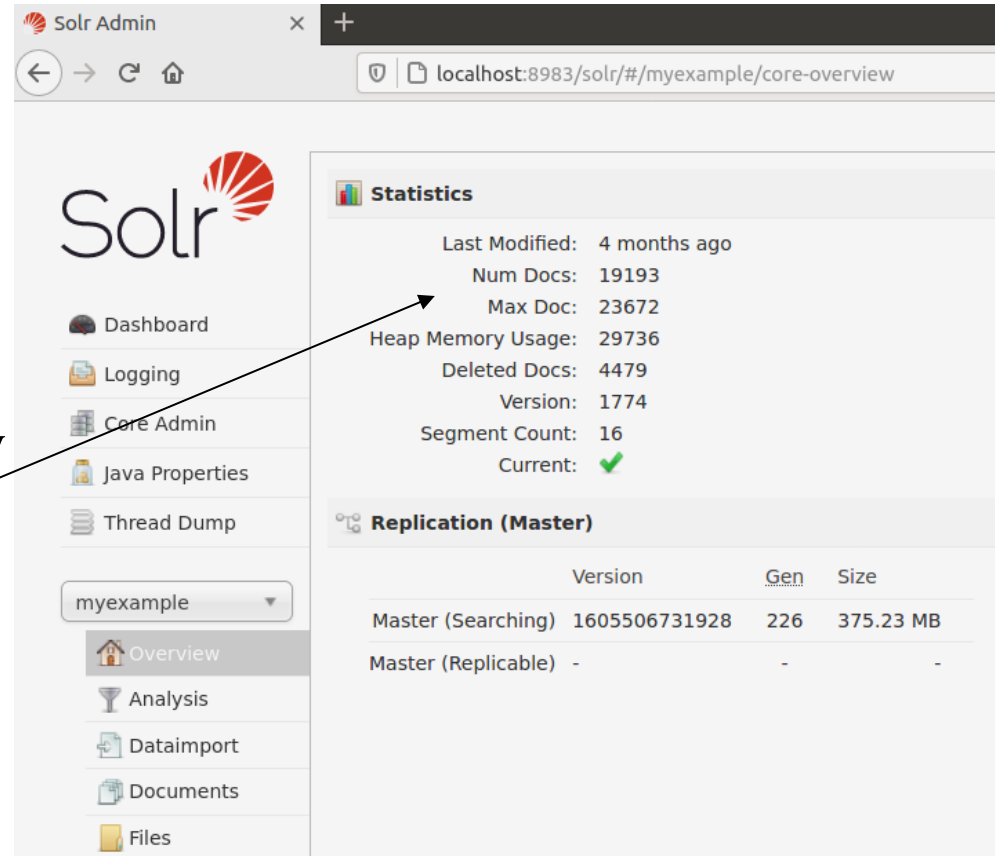
- We have created a reference set of web pages for all news websites and placed them on Google drive
- Below is a URL that will give you access to these data sets
 - Each data set contains three files
- The folders are:
 - NYTIMES [01 ~ 30]
 - FOXNEWS [31 ~ 60]
 - LATIMES [61 ~ 00]
- Download only the folder you are responsible for
 - Check your USC ID against the list in the exercise
- here is the URL again
 - <https://drive.google.com/drive/folders/14qtJywCRk00wEbmt-Pv3efyPIBqwcoZr>





Step 3: Use Solr to Index a Web Site

- start the Solr server
- start a new Solr core
- Use Tika to import your saved files
- Use the Solr interface to check that the files have been properly indexed
- Note the URL:
`localhost:8983/solr/#/myexample`
- 19193 docs successfully indexed



The screenshot shows the Solr Admin web interface in a browser. The address bar displays `localhost:8983/solr/#/myexample/core-overview`. The left sidebar contains navigation links: Dashboard, Logging, Core Admin, Java Properties, Thread Dump, and a dropdown menu for 'myexample' which is currently expanded to show Overview, Analysis, Dataimport, Documents, and Files. The main content area displays the Solr logo and a 'Statistics' section with the following data:

Statistic	Value
Last Modified:	4 months ago
Num Docs:	19193
Max Doc:	23672
Heap Memory Usage:	29736
Deleted Docs:	4479
Version:	1774
Segment Count:	16
Current:	✓

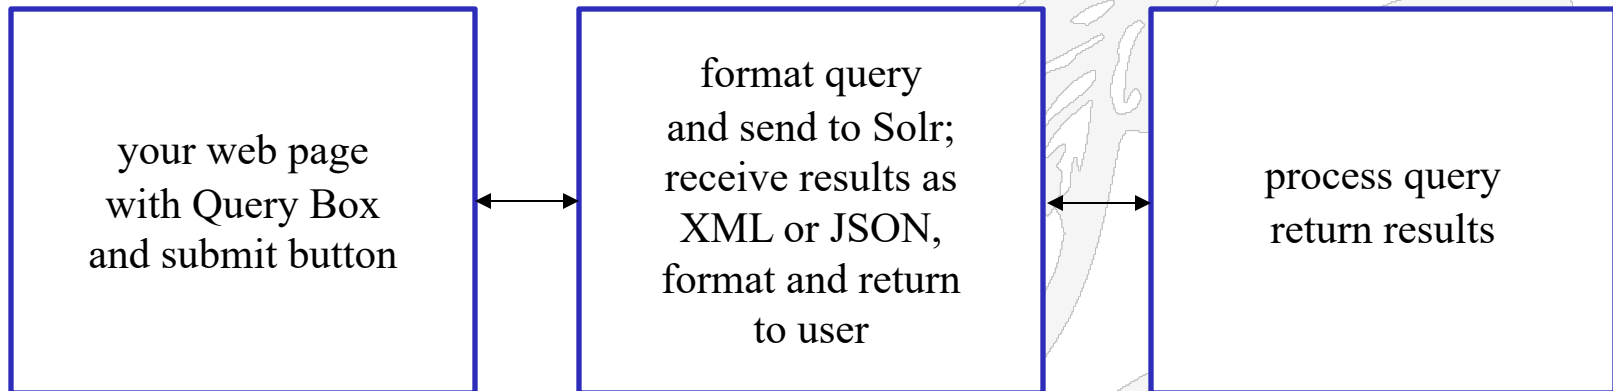
Below the statistics is a 'Replication (Master)' section with a table showing the status of the master and replicable nodes:

	Version	Gen	Size
Master (Searching)	1605506731928	226	375.23 MB
Master (Replicable)	-	-	-

An arrow from the text '19193 docs successfully indexed' in the list points to the 'Num Docs: 19193' value in the Solr Admin interface.

Step 4: Comparing Search Engine Ranking Algorithms

1. You should download the reference files for the news website you are responsible for;
2. You should install Solr as described previously;
3. Take the web pages from the reference files and index them in Solr, as described in earlier slides
4. Build a front end to Solr that permits a visitor to enter a query and get matching results
5. Solr will return the results in JSON format; your server needs to take the results and format them for the user



Web Browser

Your Web Server
Copyright Ellis Horowitz, 2011-2022Your version of Solr₂₈



Apache Solr Client APIs

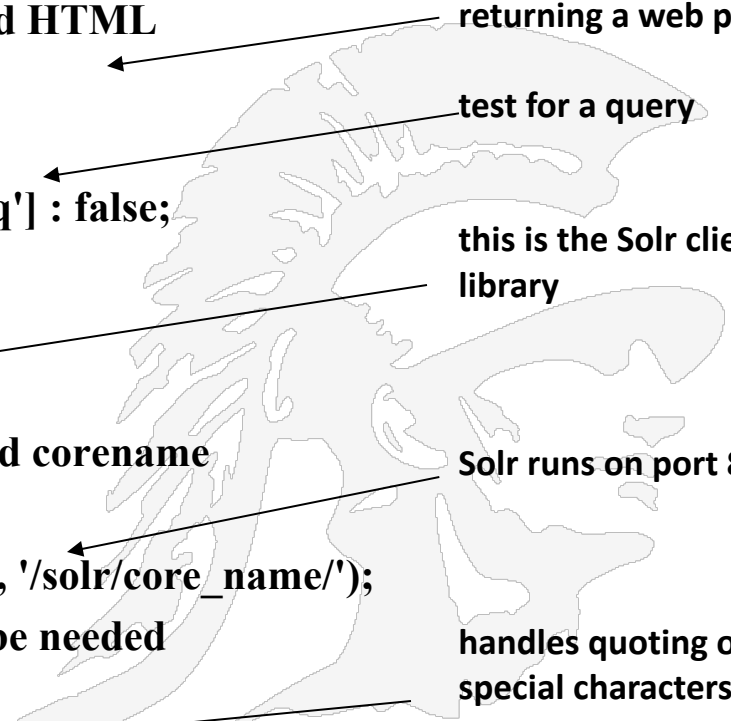
- **http is the protocol to be used between client applications and Solr**
- **Clients use Solr's five basic operations: query, index, delete, commit, and optimize**
- **JavaScript is the standard client API and there is no package to be installed**
 - http requests can be sent using XMLHttpRequest
 - Solr will respond with JSON output
- **there is an output format specifically for Python**
 - See https://lucene.apache.org/solr/guide/7_7/using-python.html
- **Full list of Solr API clients can be found here**
 - https://lucene.apache.org/solr/guide/7_7/client-api-lineup.html



Step 4

Here is a PHP client that accepts input from the user in a HTML form, and sends the request to the Solr server. After the Solr server processes the query, it returns the results which is parsed by the PHP program and formatted for display

```
<?php
// make sure browsers see this page as utf-8 encoded HTML
header('Content-Type: text/html; charset=utf-8');
$limit = 10;
$query = isset($_REQUEST['q']) ? $_REQUEST['q'] : false;
$results = false;
if ($query)
{ require_once('Apache/Solr/Service.php');
  // create a new solr service instance - host, port, and corename
  // path (all defaults in this example)
  $solr = new Apache_Solr_Service('localhost', 8983, '/solr/core_name/');
  // if magic quotes is enabled then stripslashes will be needed
  if (get_magic_quotes_gpc() == 1)
  { $query = stripslashes($query); }
}
```



returning a web page

test for a query

this is the Solr client library

Solr runs on port 8983

handles quoting of special characters in query



PHP Program (2 of 3)

```

try
{ $results = $solr->search($query, 0, $limit); }
catch (Exception $e)
{ die("<html><head><title>SEARCH EXCEPTION</title><body><pre>{$e-
>__toString()}</pre></body></html>"); } }
?>

<html> <head> <title>PHP Solr Client Example</title> </head> <body>
<form accept-charset="utf-8" method="get">
<label for="q">Search:</label>
<input id="q" name="q" type="text" value="<?php echo htmlspecialchars($query, ENT_QUOTES,
'utf-8'); ?>"/>
<input type="submit"/>
</form>
<?php
// display results
if ($results)
{ $total = (int) $results->response->numFound; $start = min(1, $total); $end = min($limit, $total); ?>
<div>Results <?php echo $start; ?> - <?php echo $end; ?> of <?php echo $total; ?>:</div>
<ol>

```

send query to Solr
 catch any exception

create web page
 output
 create input text box

create submit button

end form

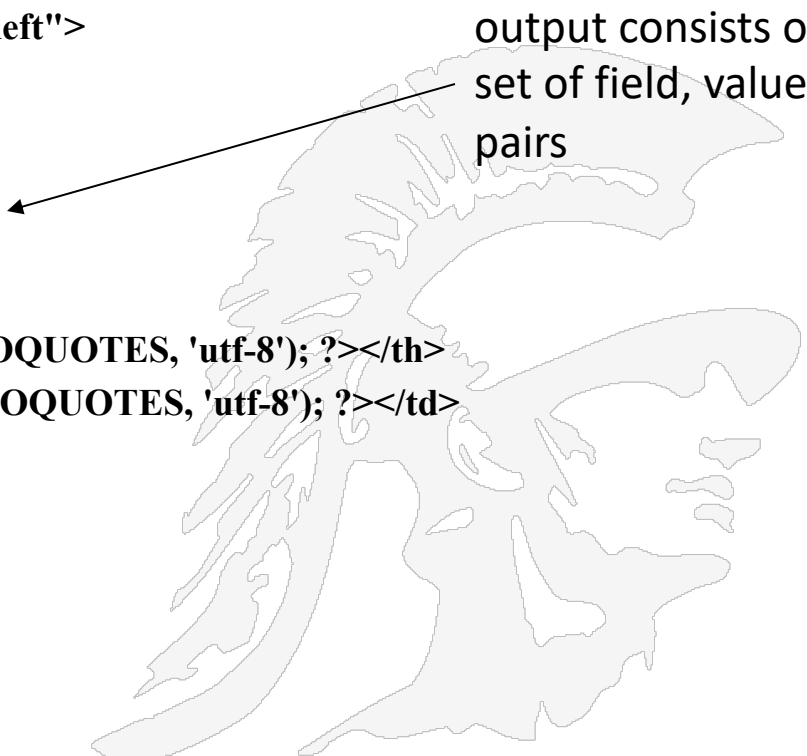
JSON result string



PhP Program (3 of 3)

```
<?php
// iterate result documents
foreach ($results->response->docs as $doc)
{ ?> <li>
  <table style="border: 1px solid black; text-align: left">
<?php
// iterate document fields / values
foreach ($doc as $field => $value)
{ ?>
  <tr>
<th><?php echo htmlspecialchars($field, ENT_NOQUOTES, 'utf-8'); ?></th>
<td><?php echo htmlspecialchars($value, ENT_NOQUOTES, 'utf-8'); ?></td>
</tr>
<?php }
?> </table> </li>
<?php } ?>
</ol>
<?php
} ?> </body> </html>
```

output consists of a
set of field, value
pairs

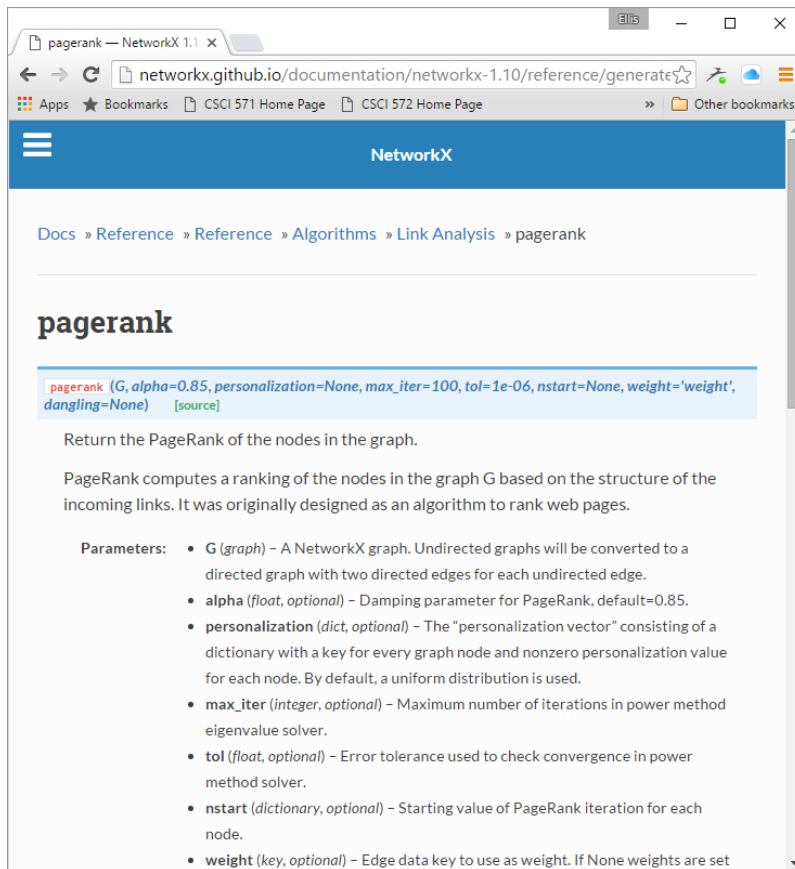




Comparing Ranking Algorithms

- **we have already seen the built-in Solr ranking method**
 - Based on Lucene
- **Solr permits alternative ranking algorithms**
 - to use PageRank we must create pageranks for all web pages in our news website
 - to do this we use an open source program called NetworkX
 - we need to create the web graph (pages and links) so the PageRank can be determined
 - See the following document for a method on extracting links and producing the NetworkX file:
 - <http://csci572.com/2022Spring/hw4/ExtractingLinks.pdf>

http://networkx.github.io/documentation/networkx-1.10/reference/generated/networkx.algorithms.link_analysis.pagerank_alg.pagerank.html



The screenshot shows a web browser window displaying the NetworkX documentation for the `pagerank` algorithm. The browser's address bar shows the URL `networkx.github.io/documentation/networkx-1.10/reference/generated/networkx.algorithms.link_analysis.pagerank_alg.pagerank.html`. The page has a blue header with the NetworkX logo and a navigation menu. The main content area shows the `pagerank` function signature: `pagerank(G, alpha=0.85, personalization=None, max_iter=100, tol=1e-06, nstart=None, weight='weight', dangling=None)`. Below the signature, there is a brief description of the function and a list of parameters with their descriptions.

pagerank

```
pagerank(G, alpha=0.85, personalization=None, max_iter=100, tol=1e-06, nstart=None, weight='weight', dangling=None)
```

Return the PageRank of the nodes in the graph.

PageRank computes a ranking of the nodes in the graph `G` based on the structure of the incoming links. It was originally designed as an algorithm to rank web pages.

Parameters:

- `G (graph)` – A NetworkX graph. Undirected graphs will be converted to a directed graph with two directed edges for each undirected edge.
- `alpha (float, optional)` – Damping parameter for PageRank, default=0.85.
- `personalization (dict, optional)` – The "personalization vector" consisting of a dictionary with a key for every graph node and nonzero personalization value for each node. By default, a uniform distribution is used.
- `max_iter (integer, optional)` – Maximum number of iterations in power method eigenvalue solver.
- `tol (float, optional)` – Error tolerance used to check convergence in power method solver.
- `nstart (dictionary, optional)` – Starting value of PageRank iteration for each node.
- `weight (key, optional)` – Edge data key to use as weight. If None weights are set

- You are going to use an open source PageRank algorithm, networkx, located at URL above;
- networkx is a Python program
- You need to create a graph that the PageRank algorithm can work on

Important Parameters:

- a NetworkX graph
- a damping parameter (e.g. 0.85)
- maximum number of iterations
- error tolerance
- starting Page Rank value of nodes
- (see expansion on next slide)



http://networkx.github.io/documentation/networkx-1.10/_modules/networkx/algorithms/link_analysis/pagerank_alg.html#pagerank

```
def pagerank(G, alpha=0.85, personalization=None,  
            max_iter=100, tol=1.0e-6, nstart=None, weight='weight', dangling=None):
```

Return the PageRank of the nodes in the graph.

Parameters

G : *graph*: A NetworkX graph.

alpha : float, optional Damping parameter for PageRank, *default*=0.85.

personalization: dict, optional; By *default*, a uniform distribution is used.

max_iter : integer, optional Maximum number of iterations in power method eigenvalue solver.

tol : float, optional; Error tolerance used to check convergence in power method solver.

nstart : dictionary, optional Starting value of PageRank iteration for each node.

weight : key, optional; If None weights are set to 1.

Returns

pagerank : dictionary; Dictionary of nodes with PageRank as value

Examples

```
>>> G = nx.DiGraph(nx.path_graph(4))
```

```
>>> pr = nx.pagerank(G, alpha=0.9)
```



Some NetworkX Operations

- Create an empty graph with no nodes and no edges.

```
>>> import networkx as nx
```

```
>>> G=nx.Graph()
```

- add one node at a time

```
>>>G.add_node(1)
```

- add a list of nodes

```
>>>G.add_nodes_from([2,3])
```

- add one edge at a time

```
>>>G.add_edge(1,2)
```

```
>>>e=(2,3)
```

```
>>>G.add_edge(*e) #unpack edge tuple*
```

- add a list of edges

```
>>>G.add_edges_from([(1,2),(1,3)])
```





Final Steps

- **Input to the PageRank algorithm a file containing every document ID and associated with each ID, the IDs that are pointed to by links within the document ID**
- **Output from the PageRank algorithm is a file containing every document ID and its associated PageRank**
- **place this file in solr-7.7.3/server/solr/core_name, call the file external_pageRankFile.txt**
- **add the PageRank field to the schema.xml file**

```
<fieldType name="external" keyField="id" defVal="0" class="solr.ExternalFileField" />
```

```
<field name="pageRankFile" type="external" stored="false" indexed="false" />
```



The Queries

- Once both ranking algorithms are working you should input the queries below and compare the results

Nine Queries
Cannes
Congress
Democrats
Patriot Movement
Republicans
Senate
Olympics 2020
Stock
Virus